Kill My Idea

Product Requirements Document and Vendor Build Brief

Prepared for external product design and development vendors

Document purpose

This document consolidates the current product vision, feature requirements, operating principles, risk controls, and suggested rollout plan for Kill My Idea. It is intended to help a software company estimate, scope, design, and build the product without relying on scattered chat notes.

Product	Kill My Idea
Category	AI-powered SaaS web platform
Primary use case	Founder idea evaluation and next-step guidance
Working positioning	Brutal startup idea validator / founder decision engine
Primary audience	First-time founders, indie hackers, and early-stage builders
Document status	Company-ready build brief (v1-focused)

Recommended build framing: market the product as a “roast my idea” tool, but design the core product as a transparent founder decision engine that can say pursue, refine, test first, drop, or insufficient data.

- This document is based on founder assumptions and product reasoning, not on completed market validation research.
- The vendor should treat this as the build baseline and run a discovery workshop to confirm unresolved assumptions before final design and engineering estimates.

1. Executive Summary

Kill My Idea is a web application that helps founders turn messy early idea thinking into a structured decision. Instead of asking users to rely on encouragement, fragmented research, or generic AI answers, the product asks clarifying questions, analyzes observable market signals, and produces a blunt recommendation with evidence and next steps.

The product should not position itself as a predictor of startup success. It should position itself as a decision support system that helps users decide whether an idea is worth exploring, testing, refining, or building right now.

---

2. Product Vision and Promise

Vision

To become the most trusted system for helping founders decide whether a startup idea is worth exploring, testing, refining, or building.

Product promise

Before building, this product helps users decide whether their idea is worth pursuing and what they should do next.

Product personality

Harsh but constructive. The experience should feel like a ruthless investor combined with a sharp cofounder. It must never feel like false encouragement software.

3. Problem Statement

- Founders often get excited about an idea, do fragmented research, and still remain unsure what decision to make.
- Existing idea validation workflows are slow, inconsistent, and require a level of product or market judgment many first-time founders do not yet have.
- Current AI tools often produce generic encouragement instead of evidence-backed guidance.
- The cost of bad validation is delayed but severe: wasted months, poor build choices, avoidable demotivation, and opportunity cost.

4. Product Goals

Primary goal	Help founders reach a clearer and better next decision faster.
Secondary goal	Create a repeatable framework for comparing ideas and tracking refinement over time.
Trust goal	Earn trust through transparency, visible assumptions, evidence snippets, and confidence levels.

Business goal	Convert curiosity-driven traffic into paid deep analyses and repeat usage across multiple idea cycles.

5. Target Users and Roles

User segments

Segment	Why they use it	Likelihood to pay	Notes
First-time technical founder	Needs a structured way to judge whether an idea is worth pursuing.	Medium	Strong top-of-funnel segment, especially students and hackathon builders.
Indie hacker / repeat builder	Compares multiple ideas and wants faster decisions.	High	Best early monetization segment due to time sensitivity and repeat usage.
Working professional with side project	Wants to avoid wasting nights and weekends on weak ideas.	High	Strong candidate for paid credits and exports.
Casual idea explorer	Mostly seeks entertainment or reassurance.	Low	Useful for traffic but not a v1 core target.

System roles

Role	Access and permissions	Priority	Notes
Visitor	Use limited quick roast, view teaser output, create account prompt.	P0	Core acquisition role.
Registered user	Save ideas, run limited analyses, refine ideas, view history.	P0	Base product role.
Paid user	Access full reports, comparisons, exports, and deeper evidence.	P0	Revenue role.
Admin / operator	Manage prompts, plans, moderation, failed jobs, and system settings.	P1	Needed for operations and support.

6. What the Product Must Not Do

- It must not pretend to predict startup success with certainty.
- It must not issue polished verdicts for vague ideas without forcing clarification first.
- It must not become a full founder operating system in v1.
- It must not hide weak evidence behind confident language.

7. Core User Journeys

Journey A — Quick Roast to Deep Analysis

1. Visitor lands on the product from marketing or a shared roast report.
2. Visitor pastes a rough idea into the Quick Roast flow.
3. System returns a short blunt reaction, highlights major concerns, and invites deeper analysis.
4. User signs up to unlock structured clarification and full signal-backed analysis.
5. User receives a full verdict, sees evidence, and decides whether to refine, test, or stop.

Journey B — Serious Evaluation of One Idea

6. User enters idea summary, target user, and problem statement.
7. System runs clarification prompts to tighten the ICP, problem severity, alternatives, and monetization assumptions.
8. System fetches relevant market signals and normalizes them by idea category.
9. Decision engine assigns sub-scores, confidence, and a verdict.
10. Report explains reasoning and recommends the next best actions.

Journey C — Compare and Prioritize Multiple Ideas

11. User stores multiple ideas in the workspace.
12. User reruns or updates analyses as the ideas evolve.
13. System compares the ideas on demand, urgency, competition, distribution, and confidence.
14. User chooses which idea to test first based on side-by-side trade-offs rather than intuition alone.

Key workflow decisions

Decision point	System behavior	Rationale
Idea too vague	Pause verdict, request clarification, and explain what is missing.	Prevents false confidence.
Signals conflict	Lower confidence and show trade-offs explicitly.	Keeps the product honest.
Low data coverage	Return “Insufficient Data” or “Test First” rather than forcing certainty.	Protects product trust.
Idea looks weak	Allow “Drop” and show salvage options if any exist.	The product must be willing to reject ideas.

8. Product Modules Overview

V1 module map

Module	Purpose	Priority	Key outputs
Quick Roast	Fast top-of-funnel reaction and marketing hook.	P0	Blunt preview verdict and upgrade trigger
Idea Intake	Collect raw inputs and basic context.	P0	Idea record and starter context
Clarification Engine	Transform vague thinking into structured evaluation input.	P0	Clearer ICP, problem, alternatives, monetization assumptions
Signal Analysis Engine	Fetch and interpret external evidence.	P0	Signal summaries and evidence snippets
Decision Engine	Generate scores, confidence, and verdict.	P0	Pursue / Refine / Test First / Drop / Insufficient Data
Report & Next Steps	Explain the decision and recommend action.	P0	Full analysis report and validation steps
Workspace & History	Save ideas, reruns, versions, and comparisons.	P1	Idea library and iteration loop
Share & Export	Enable sharing and downloadable outputs.	P1	Public links and PDF/export assets
Admin & Ops	Support moderation, billing, prompt controls, and job monitoring.	P1	Operational controls and support tooling

9.1 Quick Roast Module

Feature	Requirement / Behavior	Pri.	Notes
Public text input	Allow a visitor to paste a rough idea with minimal friction.	**P0**	Main homepage CTA.

Feature	Requirement / Behavior	Pri.	Notes
Rapid response	Return a short and opinionated reaction in a few sections: first impression, biggest flaw, what must be clarified.	**P0**	Must feel sharp, not generic.
Teaser limits	Keep output short enough that deep analysis still feels necessary.	**P0**	Protects monetization.
Upgrade trigger	Prompt sign-up for full analysis, saved ideas, or comparisons.	**P0**	Main conversion step.
Shareable roast view	Let users share the roast result through a clean public link.	**P1**	Useful growth loop.

9.2 Idea Intake and Clarification Module

Feature	Requirement / Behavior	Pri.	Notes
Minimum required inputs	Collect idea summary, target user, and problem before serious analysis begins.	**P0**	Without these, verdict quality is too weak.
Clarification prompts	Ask follow-up questions on ICP, pain severity, current alternatives, differentiation, monetization, and distribution assumptions.	**P0**	Core trust feature.
Idea statement improvement	Rewrite the idea into a clearer structured summary before analysis.	**P0**	Should help users who think vaguely.
Vagueness detection	Explicitly tell the user when the idea is too vague to evaluate properly.	**P0**	Must block dishonest outputs.

Feature	Requirement / Behavior	Pri.	Notes
Edit / confirm step	Allow the user to approve the interpreted version before signal collection.	**P1**	Reduces misanalysis.

9.3 Signal Analysis Engine

Feature	Requirement / Behavior	Pri.	Notes
Signal sources	Use v1 sources such as Google trend patterns, search results, Reddit pain discussions, SEO demand indicators, competitor examples, reviews, pricing patterns, app presence, and relevant social chatter.	**P0**	Signal quality is central to product credibility.
Category-aware weighting	Interpret signals differently for B2B, consumer, devtool, and similar product categories.	**P0**	Signal meaning varies by category.
Evidence snippets	Store raw evidence excerpts or short summaries that support conclusions.	**P0**	Needed for explainability.
Source traceability	Link each conclusion to the signal category that influenced it.	**P0**	Users must understand why the tool said what it said.
Low-confidence handling	When signals are thin or conflicting, reduce confidence and avoid strong verdicts.	**P0**	Essential risk control.

9.4 Decision Engine

Feature	Requirement / Behavior	Pri.	Notes
Sub-scores	Calculate scores for demand, urgency, competition pressure, differentiation, distribution feasibility, monetization, execution difficulty, and confidence.	**P0**	Sub-scores are more useful than one raw number.
Overall score	Generate a single score that represents probability worth testing, not probability of success.	**P0**	Language must stay honest.
Verdict system	Return one of five states: Pursue, Refine, Test First, Drop, or Insufficient Data.	**P0**	This is the product center.
Flexible logic	Use score ranges as guidance but allow logic-based overrides when signals conflict.	**P0**	Prevents over-simplified scoring.
Visible assumptions	Display assumptions used in the interpretation and verdict.	**P0**	Protects credibility.

9.5 Report and Next-Step Module

Feature	Requirement / Behavior	Pri.	Notes
Full report sections	Include verdict, overall score, sub-scores, confidence, reasoning, assumptions, red flags, green flags, signal evidence, and next actions.	**P0**	Primary paid output.
Refinement suggestions	If the idea is weak, suggest narrower ICPs, stronger wedges, repositioning, or monetization tweaks.	**P0**	Users need more than “no”.

Feature	Requirement / Behavior	Pri.	Notes
Test plan output	For “Test First”, propose interviews, landing pages, outreach, or waitlist experiments.	**P0**	Most practical post-verdict outcome.
Pursue warning layer	Even positive verdicts must include unknowns and major risks.	**P0**	Prevents false confidence.
Founder-ready exports	Offer PDF / downloadable summary and premium downstream outputs later.	**P1**	Useful for paid conversion and sharing.

9.6 Workspace, History, and Comparison Module

Feature	Requirement / Behavior	Pri.	Notes
Saved ideas	Store all user ideas with titles, statuses, timestamps, and tags.	**P1**	Supports repeated usage.
Version history	Track how each idea changes over time and preserve previous analyses.	**P1**	Useful for refinement loops.
Compare ideas	Allow two or more ideas to be compared side by side on key dimensions.	**P1**	Strong repeat-use feature.
Rerun analysis	Support reruns after edits or new signal pulls.	**P1**	Should create product stickiness.
Workspace organization	Include simple folders, tags, or filters for idea management.	**P2**	Helpful but non-core.

9.7 Share, Export, and Billing Module

Feature	Requirement / Behavior	Pri.	Notes
Public share links	Generate safe read-only links for roast results or full reports.	**P1**	Growth and utility.
Export to PDF	Allow paid users to download polished reports.	**P1**	Vendor should design a strong export template.
Credit-based access	Charge credits for deep analysis or premium report generation.	**P0**	Best-fit monetization model for cyclical use.
Plan management	Support free, starter, and extra-credit purchases.	**P1**	Needed for monetization scaling.
Usage gating	Control which outputs are visible in teaser vs paid mode.	**P0**	Critical business logic.

9.8 Admin and Operations Module

Feature	Requirement / Behavior	Pri.	Notes
Prompt and policy controls	Manage analysis prompts, copy variants, and safety boundaries.	**P1**	Essential for iteration and QA.
Analysis job monitor	Track queued, running, completed, and failed analyses.	**P1**	Important if signal collection is asynchronous.
Billing support tools	View user credits, plan state, and manual adjustments.	**P1**	Needed for customer support.
Moderation / abuse tools	Handle spam inputs, malicious content, or public-share misuse.	**P1**	Required for a public growth loop.
Analytics dashboard	Monitor conversion, report completion, and quality feedback.	**P1**	Supports product decisions.

10. Recommended Scoring and Verdict Design

The company should implement the scoring system as configurable business logic rather than hard-coded constants. An initial recommended model is below.

Suggested starting score model

Dimension	What it measures	Suggested weight	Notes
Demand	Evidence that people care or search for the problem area.	20%	High importance.
Problem urgency	How painful or frequent the problem appears to be.	20%	High importance.
Distribution feasibility	Whether the founder can realistically reach users.	20%	Often overlooked, should be prominent.
Differentiation	Whether the product angle is sufficiently distinct.	12%	Critical in crowded spaces.
Competition pressure	How hard the existing market may be to enter.	10%	Must be interpreted carefully.
Monetization potential	Whether pricing or willingness to pay appears plausible.	10%	Should not dominate v1 verdicts.
Execution difficulty	How hard it may be to build, launch, and sustain.	8%	Useful but secondary to demand and urgency.

Verdict logic guardrails

- A high overall score should not automatically mean “Pursue” if distribution is unrealistic or the idea remains unclear.
- A low overall score should not always mean “Drop” if the concept shows promise but needs narrower positioning.
- “Insufficient Data” should be used when evidence or input quality is too weak for a responsible recommendation.
- “Test First” should be treated as a valid and positive outcome, not as a disguised rejection.

11. Full Report Structure

- Idea summary and structured interpretation
- Verdict and overall score
- Sub-scores and confidence level
- Reasoning and assumptions used
- Demand signals and competition observations

- Differentiation gaps and monetization comments
- Red flags and green flags
- What makes the idea weak in its current form
- What could make the idea stronger
- Recommended next step and validation actions

12. Visual Design Requirements for the Report

- Verdict badge with strong hierarchy
- Sub-score cards or compact score table
- Red / yellow / green flags presented in a scannable format
- Confidence indicator separate from the verdict
- Clear separation between evidence, interpretation, and action recommendations
- Professional export design that is useful for founders and shareable with collaborators or mentors

13. Monetization and Plan Design

Commercial model recommendation

Item	Recommendation	Why
Primary pricing shape	Freemium plus credit-based deep analysis	Matches cyclical founder usage better than pure monthly subscription.
Free offering	Limited quick roasts, teaser outputs, and restricted storage	Supports acquisition and sharing.
Paid core	Full reports, evidence traceability, comparisons, and exports	Directly tied to perceived value.
Premium later	Validation experiment generator, GTM brief, PRD draft	Creates monetizable expansion paths without bloating v1.

14. High-Level Data Model

- User — account, profile, plan, credit balance, status, preferences
- Idea — title, raw idea text, target user, problem statement, status, owner, tags
- IdeaVersion — structured interpretation, clarification answers, timestamp, editor context
- AnalysisRun — inputs snapshot, signal fetch status, scoring output, verdict, confidence, model metadata
- SignalEvidence — source type, evidence snippet, normalized summary, strength, URL or trace reference
- ReportExport — export type, generated file, share settings, created timestamp
- CreditTransaction — purchase, deduction, refund, adjustment, related analysis run
- ShareLink — target report, visibility, expiry, analytics

15. Non-Functional Requirements

Operational and quality requirements

Area	Requirement	Priority	Notes
Performance	Quick Roast should feel fast; deep analysis may run as a managed background job with clear progress states.	P0	Avoid a “frozen” UX during signal fetches.
Reliability	Analysis jobs must be retryable and observable. Failures should surface a clean status and next action.	P0	Critical for trust.
Explainability	Every meaningful conclusion should map back to evidence or user-supplied assumptions.	P0	Core product principle.
Security	Protect user accounts, reports, billing data, and shared links. Use secure auth and role-based access.	P0	Baseline SaaS requirement.
Privacy	Only store the data required to support the workspace and report history; be clear about evidence storage.	P0	Important for user trust and compliance.
Scalability	Architecture should support asynchronous signal collection and report generation.	P1	Helps with expansion.
Analytics	Track conversion, rerun rate, comparison usage, share usage, and verdict distributions.	P1	Needed for product learning.
Admin control	Operators should be able to disable signal sources or prompt versions without redeploying the full application.	P1	Useful for live iteration.

16. Key Risks and Mitigations

Primary product risks

Risk	Why it matters	Mitigation
False certainty	Users may over-trust verdicts and treat them as guarantees.	Make confidence, assumptions, and low-data states highly visible.
Weak input quality	Bad idea descriptions lead to bad recommendations.	Force clarification before serious analysis.
Generic outputs	The product may feel like dressed-up prompting.	Use structured scoring, evidence traceability, and category-aware analysis.
Signal noise	External signals can conflict or mislead.	Show trade-offs, weight signals by category, and reduce confidence when needed.
Low retention	Users may only visit once.	Build saved ideas, reruns, compare mode, and validation next-step utilities.
Gimmick branding	“Kill My Idea” may attract clicks but reduce seriousness.	Use the bold brand for acquisition but keep the in-product experience professional and transparent.

17. Phased Delivery Recommendation

Suggested rollout sequence

Phase	Scope	Success condition
Phase 1	Landing experience, Quick Roast, sign-up, core idea intake, clarification engine, basic signal pipeline, verdict, and full report.	Users can complete one honest end-to-end analysis.
Phase 2	Workspace, saved ideas, reruns, idea comparison, PDF export, public share links, basic billing.	Product supports repeat usage and monetization.
Phase 3	Improved evidence explorer, validation experiment generator, premium downstream outputs such as GTM brief or PRD draft.	Product starts expanding beyond pure validation.

18. Explicitly Out of Scope for Early Versions

- Full project management or task execution tooling
- No-code app or landing-page builder
- Large collaboration suite
- Founder social community or public idea feed
- Investor matching platform
- Custom research consulting workflow as a system requirement

19. Vendor Discovery Questions to Resolve Before Final Estimation

- Which signal sources are legally and technically acceptable in v1, and what rate limits or scraping constraints apply?
- What geographic markets should the product optimize for first, especially when interpreting demand or social signals?
- Should the product launch with only English support or plan for localization-ready architecture from day one?
- What depth of evidence storage is acceptable from a privacy, cost, and UX perspective?
- Should deep analysis run synchronously for small jobs or as an explicit queued job with progress updates?
- What level of export quality is needed in v1 — basic PDF, branded PDF, or editable document outputs later?
- What should the first billing package look like in actual credit numbers and report limits?
- How much manual admin override is needed for support, abuse handling, prompt changes, and billing adjustments?

20. Final Recommendation

The product should be sold with a sharp hook but built with sober product discipline. “Roast my idea” is a strong acquisition message. The actual software should be an explainable founder decision engine that is willing to say no, willing to say “test first”, and equally willing to admit when the evidence is weak.

For a software company building this product, the highest leverage is not flashy AI output. It is the combination of clarity enforcement, transparent evidence handling, configurable verdict logic, and a user experience that makes founders feel they are getting judgment rather than fluff.

Recommended product verdict on the concept itself: refine the framing, then pursue the build.